

```
In 1 1 # So we will combine the ideas from the losses we have experimented with
2 # InverseDepth Loss works fine if there is no feathering in the subject . But if we penalize also on incorrectly detecting edges . we might
... improve upon the loss . So a edge-guide is needed
3
4 # And we will linearly combine with the loss that we experimented with Online sampling
Executed at 2023.12.10 18:16:31 in 4ms
```

```
In 4 1 import torch
2 from torch import nn
3 import numpy as np
4 import torch.nn.functional as F
5 from torchvision import transforms
6
7 from data.loaders.DataLoader import RedWebDataset , Rescale , RandomCrop
8 from torch.utils.data import DataLoader
9
10 normal_dataset = RedWebDataset(root_dir='../data/RedWeb_V1',transform=transforms.Compose([
11     Rescale((256,256)),
12 ]))
13 batcher = DataLoader(normal_dataset,batch_size=1,shuffle=True)
Executed at 2023.12.10 18:21:34 in 9ms
```

```
In 6 1 # Lets Device our online sampling first
2
3 # Note this is derived completely from the Redweb paper
4 def onlineSampling(inputs, targets, masks, threshold, sample_num):
5
6     # find A-B point pairs from predictions (mostly random)
7     inputs_index = torch.masked_select(inputs, targets.gt(threshold))
8     num_effect_pixels = len(inputs_index)
9     shuffle_effect_pixels = torch.randperm(num_effect_pixels).cuda()
10    rgb_a = inputs_index[shuffle_effect_pixels[0:sample_num*2:2]]
11    rgb_b = inputs_index[shuffle_effect_pixels[1:sample_num*2:2]]
12
13    # find corresponding pairs from ground truth
14    depth_index = torch.masked_select(targets, targets.gt(threshold))
15    depth_a = depth_index[shuffle_effect_pixels[0:sample_num*2:2]]
16    depth_b = depth_index[shuffle_effect_pixels[1:sample_num*2:2]]
17
18    # only compute the losses of point pairs with valid ground truth i.e consistent masked
19    consistent_masks_index = torch.masked_select(masks, targets.gt(threshold))
20    consistent_masks_A = consistent_masks_index[shuffle_effect_pixels[0:sample_num*2:2]]
21    consistent_masks_B = consistent_masks_index[shuffle_effect_pixels[1:sample_num*2:2]]
22
23    # The amount of A and B should be the same!!
24    if len(depth_a) > len(depth_b):
25        depth_a = depth_a[:-1]
26        rgb_a = rgb_a[:-1]
27        consistent_masks_A = consistent_masks_A[:-1]
28
29    return rgb_a, rgb_b, depth_a, depth_b, consistent_masks_A, consistent_masks_B
Executed at 2023.12.10 18:32:20 in 2ms
```

```
In 7 1 # now lets penalize wrong edges
2 # i.e if the edges derived from the depth map does not map with the edges derived from the original image. there should be a corresponding
... penalty foir it.
3
4 # convenience wrapper function to get pixels
5 def ind2sub(idx, cols):
6     r = idx / cols
7     c = idx - r * cols
8     return r, c
9
10
11 def sub2ind(r, c, cols):
12     idx = r * cols + c
13     return idx
14
15 Executed at 2023.12.10 19:26:27 in 5ms
```

```
In 8 1 def edgeGuidedSampling(inputs, targets, edges_img, thetas_img, masks, h, w):
2
3     # find edges
4     edges_max = edges_img.max()
5     edges_mask = edges_img.ge(edges_max*0.1)
6     edges_loc = edges_mask.nonzero()
7
8     inputs_edge = torch.masked_select(inputs, edges_mask)
9     targets_edge = torch.masked_select(targets, edges_mask)
10    thetas_edge = torch.masked_select(thetas_img, edges_mask)
11    minlen = inputs_edge.size()[0]
12
13    # find anchor points (i.e, edge points)
14    sample_num = minlen
15    index_anchors = torch.randint(
16        0, minlen, (sample_num,), dtype=torch.long).cuda()
17    anchors = torch.gather(inputs_edge, 0, index_anchors)
18    theta_anchors = torch.gather(thetas_edge, 0, index_anchors)
19    row_anchors, col_anchors = ind2sub(edges_loc[index_anchors].squeeze(1), w)
20    # compute the coordinates of 4-points, distances are from [2, 30]
21    distance_matrix = torch.randint(2, 31, (4, sample_num)).cuda()
22    pos_or_neg = torch.ones(4, sample_num).cuda()
23    pos_or_neg[:,2, :] = -pos_or_neg[:,2, :]
24    distance_matrix = distance_matrix.float() * pos_or_neg
25    col = col_anchors.unsqueeze(0).expand(4, sample_num).long(
26    ) + torch.round(distance_matrix.double() * torch.cos(theta_anchors).unsqueeze(0)).long()
27    row = row_anchors.unsqueeze(0).expand(4, sample_num).long(
28    ) + torch.round(distance_matrix.double() * torch.sin(theta_anchors).unsqueeze(0)).long()
29
30    # constrain  $\theta < c \leq w$ ,  $\theta \leq r \leq h$ 
31    # Note: index should minus 1
32    col[col < 0] = 0
33    col[col > w-1] = w-1
34    row[row < 0] = 0
35    row[row > h-1] = h-1
36
37    # a-b, b-c, c-d
38    a = sub2ind(row[0, :], col[0, :], w)
39    b = sub2ind(row[1, :], col[1, :], w)
40    c = sub2ind(row[2, :], col[2, :], w)
41    d = sub2ind(row[3, :], col[3, :], w)
42    A = torch.cat((a, b, c), 0)
43    B = torch.cat((b, c, d), 0)
44
45    rgb_a = torch.gather(inputs, 0, A.long())
46    rgb_b = torch.gather(inputs, 0, B.long())
47    depth_a = torch.gather(targets, 0, A.long())
48    depth_b = torch.gather(targets, 0, B.long())
49    masks_A = torch.gather(masks, 0, A.long())
50    masks_B = torch.gather(masks, 0, B.long())
51
52    return rgb_a, rgb_b, depth_a, depth_b, masks_A, masks_B, sample_num
Executed at 2023.12.10 19:26:39 in 13ms
```

```
In _ 1 class EdgeguidedRankingLoss(nn.Module):
2
3     def __init__(self, point_pairs=10000, sigma=0.03, alpha=1.0, mask_value=-1e-8):
4         super(EdgeguidedRankingLoss, self).__init__()
5         self.point_pairs = point_pairs # number of point pairs
6         self.sigma = sigma # used for determining the ordinal relationship between a selected pair
7         self.alpha = alpha # used for balancing the effect of = and (<,>)
8         self.mask_value = mask_value
9         # self.regularization_loss = GradientLoss(scales=4)
10
11    def getEdge(self, images):
12        n, c, h, w = images.size()
13        a = torch.Tensor([[-1, 0, 1], [-2, 0, 2], [-1, 0, 1]]
14        ).cuda().view((1, 1, 3, 3)).repeat(1, 1, 1, 1)
15        b = torch.Tensor([[1, 2, 1], [0, 0, 0], [-1, -2, -1]]
16        ).cuda().view((1, 1, 3, 3)).repeat(1, 1, 1, 1)
17        if c == 3:
18            gradient_x = F.conv2d(images[:, 0, :, :].unsqueeze(1), a)
19            gradient_y = F.conv2d(images[:, 0, :, :].unsqueeze(1), b)
20        else:
21            gradient_x = F.conv2d(images, a)
22            gradient_y = F.conv2d(images, b)
23        edges = torch.sqrt(torch.pow(gradient_x, 2) + torch.pow(gradient_y, 2))
24        edges = F.pad(edges, (1, 1, 1, 1), "constant", 0)
25        thetas = torch.atan2(gradient_y, gradient_x)
26        thetas = F.pad(thetas, (1, 1, 1, 1), "constant", 0)
27        return edges, thetas
28
29    def forward(self, inputs, targets, images, masks=None):
30        if masks == None:
31            masks = targets > self.mask_value
32            # Comment this line if you don't want to use the multi-scale gradient matching term !!!
33            # regularization_loss = self.regularization_loss(inputs.squeeze(1), targets.squeeze(1), masks.squeeze(1))
34            # find edges from RGB
35            edges_img, thetas_img = self.getEdge(images)
36
37            # =====
38            n, c, h, w = targets.size()
39            if n != 1:
40                inputs = inputs.view(n, -1).double()
41                targets = targets.view(n, -1).double()
42                masks = masks.view(n, -1).double()
43                edges_img = edges_img.view(n, -1).double()
44                thetas_img = thetas_img.view(n, -1).double()
45            else:
46                inputs = inputs.contiguous().view(1, -1).double()
47                targets = targets.contiguous().view(1, -1).double()
48                masks = masks.contiguous().view(1, -1).double()
49                edges_img = edges_img.contiguous().view(1, -1).double()
50                thetas_img = thetas_img.contiguous().view(1, -1).double()
51
52            # initialization
53            loss = torch.DoubleTensor([0.0]).cuda()
54
55            for i in range(n):
56                # Edge-Guided sampling
57                rgb_a, rgb_b, depth_a, depth_b, masks_A, masks_B, sample_num = edgeGuidedSampling(
58                    inputs[i, :], targets[i, :], edges_img[i], thetas_img[i], masks[i, :], h, w)
59                # Random Sampling
60                random_sample_num = sample_num
61                random_rgb_a, random_rgb_b, random_depth_a, random_depth_b, random_masks_A, random_masks_B = onlineSampling(
62                    inputs[i, :], targets[i, :], masks[i, :], self.mask_value, random_sample_num)
63
64                # Combine EGS + RS
65                rgb_a = torch.cat((rgb_a, random_rgb_a), 0)
66                rgb_b = torch.cat((rgb_b, random_rgb_b), 0)
67                depth_a = torch.cat((depth_a, random_depth_a), 0)
68                depth_b = torch.cat((depth_b, random_depth_b), 0)
69                masks_A = torch.cat((masks_A, random_masks_A), 0)
70                masks_B = torch.cat((masks_B, random_masks_B), 0)
71
72                # GT ordinal relationship
73                target_ratio = torch.div(depth_a+1e-6, depth_b+1e-6)
74                mask_eq = target_ratio.lt(
75                    1.0 + self.sigma) * target_ratio.gt(1.0/(1.0+self.sigma))
76                labels = torch.zeros_like(target_ratio)
77                labels[target_ratio.ge(1.0 + self.sigma)] = 1
78                labels[target_ratio.le(1.0/(1.0+self.sigma))] = -1
79
80                # consider forward-backward consistency checking, i.e, only compute losses of point pairs with valid GT
81                consistency_mask = masks_A * masks_B
82
83                equal_loss = (rgb_a - rgb_b).pow(2) * \
84                    mask_eq.double() * consistency_mask
85                unequal_loss = torch.log(
86                    1 + torch.exp((-rgb_a + rgb_b) * labels)) * (~mask_eq).double() * consistency_mask
87
88                # Please comment the regularization term if you don't want to use the multi-scale gradient matching loss !!!
89                # + 0.2 * regularization_loss.double()
90                loss = loss + self.alpha * equal_loss.mean() + 1.0 * unequal_loss.mean()
91
92            return loss[0].float()/n
```